GENAI HANDSON

Name: Ananya Reddy Gandlaparthi

SRN: PES2UG23CS064

Section: A

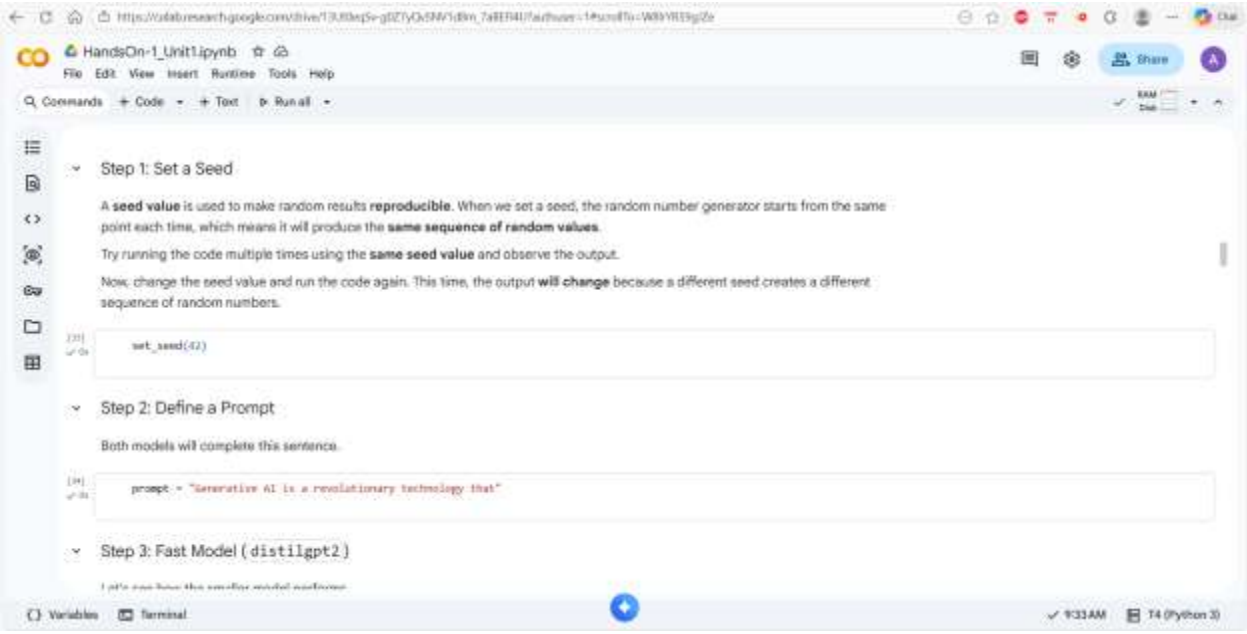

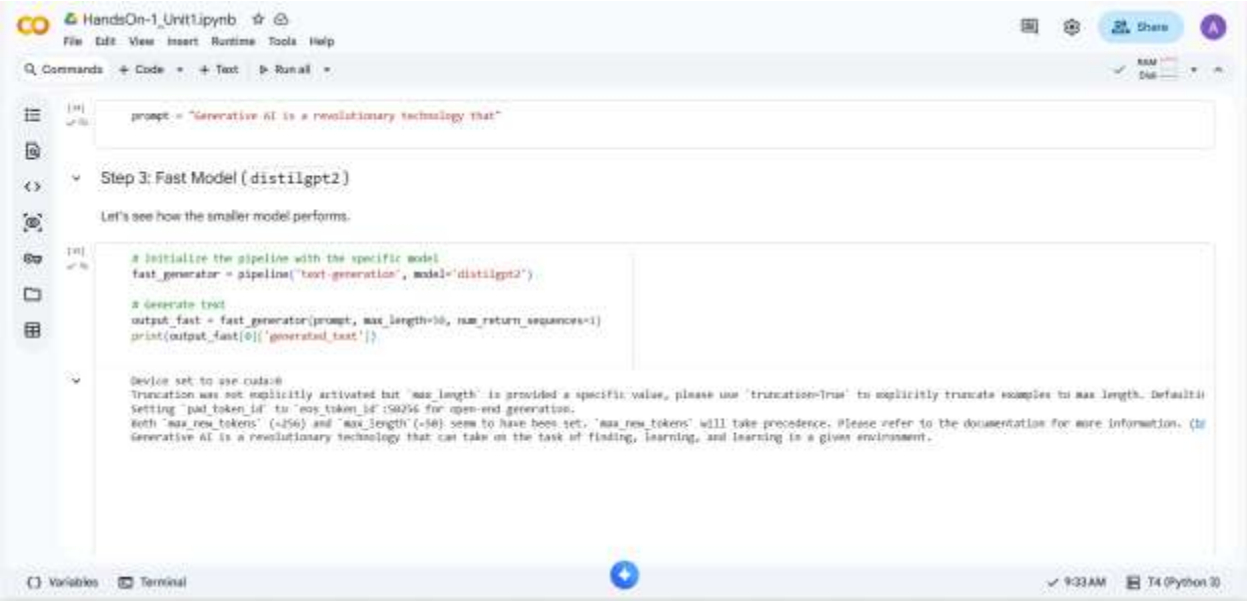

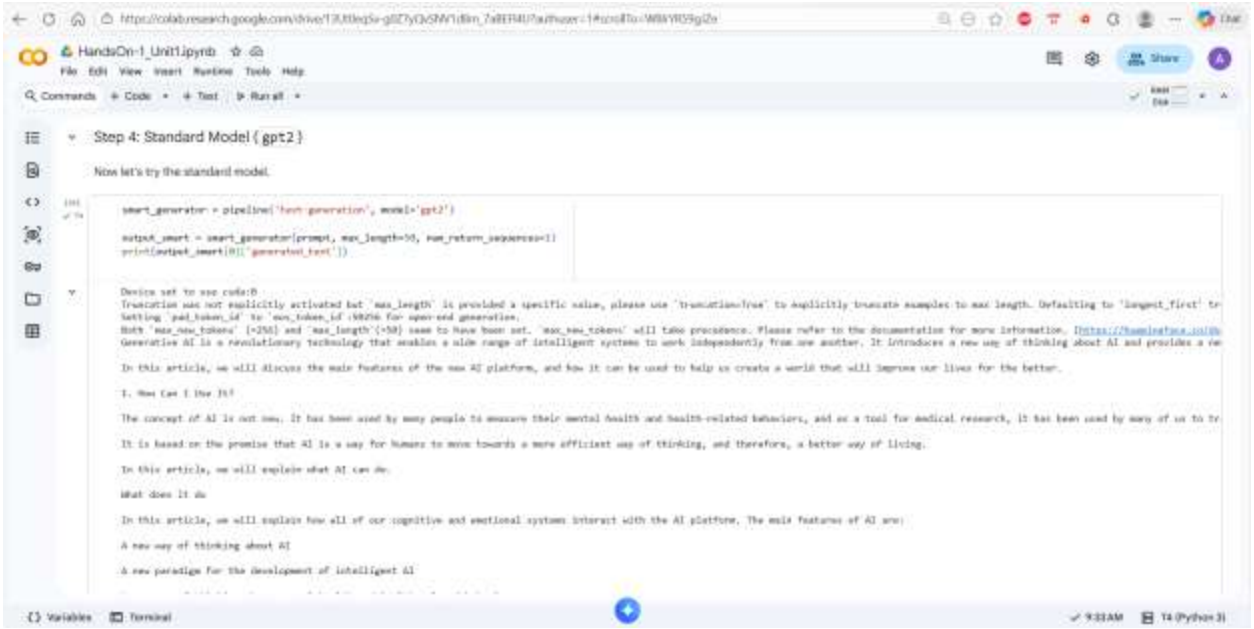
HandsOn-1_Unit1.ipynb
File Edit View Insert Runtime Tools Help
Share
Step 4: Standard Model (gpt2)
Now let's try the standard model.
smart_generator = pipeline('text-generation', model='gpt2')
output_smart = smart_generator(prompt, max_length=50, num_return_sequences=1)
print(output_smart[0]['generated_text'])
Variables
Terminal
T4 (Python 3)

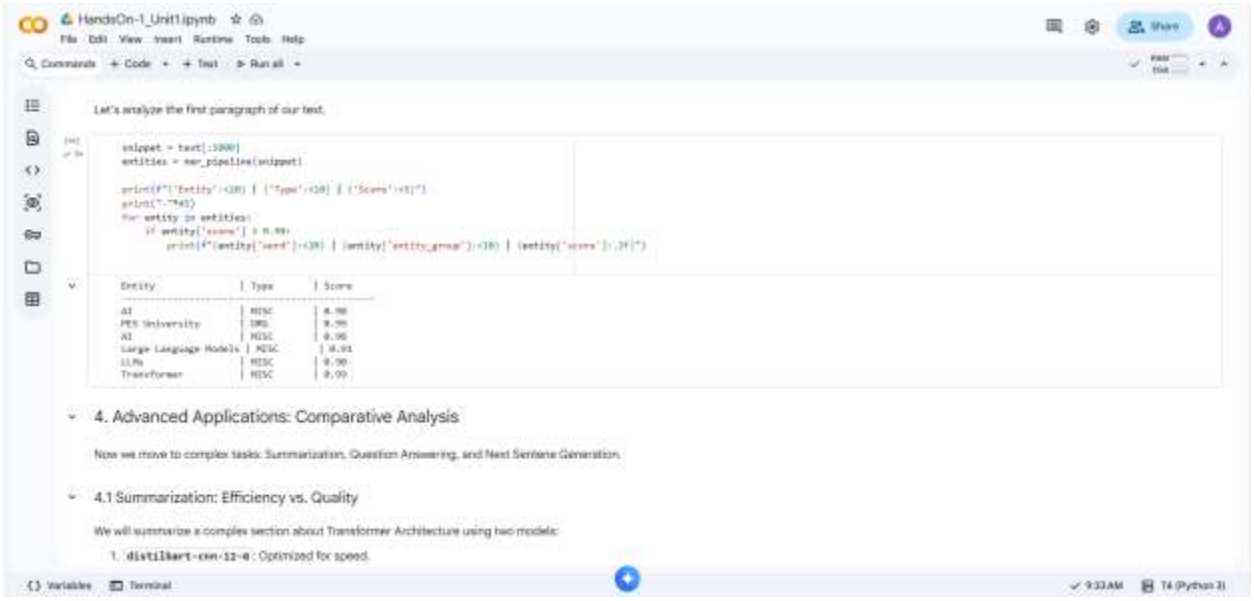
HandsOn-1_Unit1.ipynb
File Edit View Insert Runtime Tools Help
Share
Let's analyze the first paragraph of our text.
snippet = text[:1000]
entities = ner_pipeline(snippet)
Entity | Type | Score
AI | MISC | 0.98
PES University | ORG | 0.99
AI | MISC | 0.96
Large Language Models | MISC | 0.91
LLMs | MISC | 0.90
Transformer | MISC | 0.99
4. Advanced Applications: Comparative Analysis
Now we move to complex tasks: Summarization, Question Answering, and Next Sentence Generation.
4.1 Summarization: Efficiency vs. Quality
We will summarize a complex section about Transformer Architecture using two models:
1. distilbart-cnn-12-6 : Optimized for speed.
Variables
Terminal
T4 (Python 3)

Let's ask about the risks mentioned in our text.

```
questions = [
    "What is the fundamental innovation of the Transformer?",
    "What are the risks of using Generative AI?"
]

for q in questions:
    res = qa_pipeline(question=q, context=text[:1000])
    print(f"\nQ: {q}")
    print(f"A: {res['answer']}")
```

```
Q: What is the fundamental innovation of the Transformer?
A: to identify hidden patterns, structures, and relationships within the data

Q: What are the risks of using Generative AI?
A: data privacy, intellectual property, and academic integrity
```

Let's see what the model thinks Generative AI creates.

```
masked_sentence = "The goal of Generative AI is to create new [MASK]."
preds = mask_filler(masked_sentence)

for p in preds:
    print(f"{p['token_str']}: {p['score']:.2f}")
```

```
applications: 0.06
ideas: 0.05
problems: 0.05
systems: 0.04
information: 0.03
```

Seed: 50

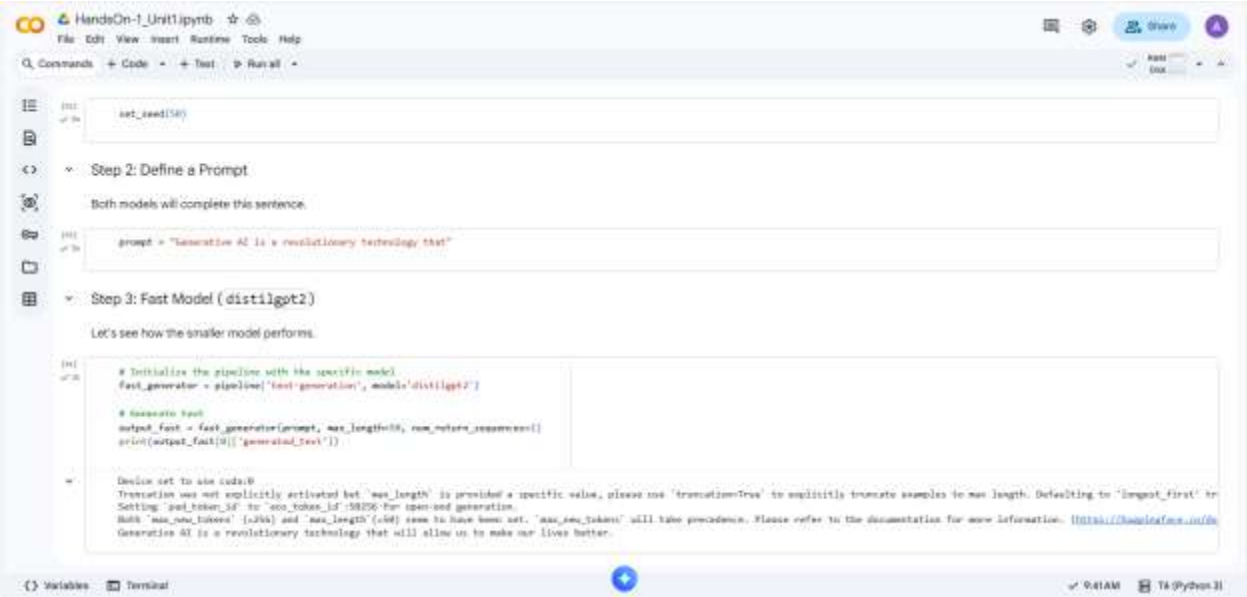

The distilgpt2 model, which is smaller and faster, generated a short, simple sentence. It was quick and gave a basic answer, but the output was limited. GPT2 model, which is bigger and smarter, generated much longer text with more details about AI benefits. It stayed more on topic and provided more information. The seed value of 42 made sure that both runs produced the same output each time.

Using the same seed value of 42 gave me the same generated text every time I ran the code, but when I changed the seed to 50 I got a different output. Bigger models like gpt2 and bart-large produce much better quality outputs than smaller models like distilgpt2 and distilbart, but they are slower and require more computing power.